

DSC680 Project3: Breast Cancer Survival Prediction

Assignment 10.1

Name: Madhuri Basava

Date: 08/02/2024

```
In [1]: # pip install xgboost imbalanced-learn scikit-optimize sweetviz
```

```
In [2]: # Import the necessary Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, OrdinalEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve, roc_auc_score, confusion_matrix, accuracy_score
import sweetviz as sv
from sklearn.model_selection import train_test_split
from xgboost import XGBClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.preprocessing import LabelEncoder
from sklearn.feature_selection import mutual_info_classif
from scipy import stats as ss
from scipy.stats import chi2_contingency
from sklearn.preprocessing import OneHotEncoder
from imblearn.over_sampling import SMOTE

# Ignore warnings
import warnings
warnings.filterwarnings('ignore')

%matplotlib inline
```

```
In [3]: # Load the Breast Cancer dataset into the data frame
breast_cancer_df = pd.read_csv('BRCA.csv')
breast_cancer_df
```

Out[3]:

	Patient_ID	Age	Gender	Protein1	Protein2	Protein3	Protein4	Tumour_Stage	Histology
0	TCGA-D8-A1XD	36.0	FEMALE	0.080353	0.42638	0.54715	0.273680	III	Infiltrating Ductal Carcinoma
1	TCGA-EW-A1OX	43.0	FEMALE	-0.420320	0.57807	0.61447	-0.031505	II	Mucinous Carcinoma
2	TCGA-A8-A079	69.0	FEMALE	0.213980	1.31140	-0.32747	-0.234260	III	Infiltrating Ductal Carcinoma
3	TCGA-D8-A1XR	56.0	FEMALE	0.345090	-0.21147	-0.19304	0.124270	II	Infiltrating Ductal Carcinoma
4	TCGA-BH-A0BF	56.0	FEMALE	0.221550	1.90680	0.52045	-0.311990	II	Infiltrating Ductal Carcinoma
...
336	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
337	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
338	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
339	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
340	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

341 rows × 10 columns



```
In [4]: # Display the information to understand the dataset
breast_cancer_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 341 entries, 0 to 340
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Patient_ID            334 non-null    object
 1   Age                   334 non-null    float64
 2   Gender                334 non-null    object
 3   Protein1              334 non-null    float64
 4   Protein2              334 non-null    float64
 5   Protein3              334 non-null    float64
 6   Protein4              334 non-null    float64
 7   Tumour_Stage          334 non-null    object
 8   Histology             334 non-null    object
 9   ER status             334 non-null    object
10   PR status             334 non-null    object
11   HER2 status           334 non-null    object
12   Surgery_type          334 non-null    object
13   Date_of_Surgery       334 non-null    object
14   Date_of_Last_Visit    317 non-null    object
15   Patient_Status        321 non-null    object
dtypes: float64(5), object(11)
memory usage: 42.8+ KB

```

```

In [5]: # Check for any missing values
breast_cancer_df.isna().sum()

```

```

Out[5]: Patient_ID            7
Age                        7
Gender                     7
Protein1                   7
Protein2                   7
Protein3                   7
Protein4                   7
Tumour_Stage               7
Histology                  7
ER status                  7
PR status                  7
HER2 status                7
Surgery_type               7
Date_of_Surgery            7
Date_of_Last_Visit        24
Patient_Status             20
dtype: int64

```

```

In [6]: # Drop the missing values
# breast_cancer_df.dropna(how='all', inplace=True)
# breast_cancer_df.dropna(subset='Patient_Status', inplace = True)

# Filter rows with at least one NaN value
breast_cancer_df = breast_cancer_df.dropna()
breast_cancer_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Index: 317 entries, 0 to 333
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Patient_ID            317 non-null    object
1   Age                   317 non-null    float64
2   Gender                317 non-null    object
3   Protein1              317 non-null    float64
4   Protein2              317 non-null    float64
5   Protein3              317 non-null    float64
6   Protein4              317 non-null    float64
7   Tumour_Stage          317 non-null    object
8   Histology              317 non-null    object
9   ER status             317 non-null    object
10  PR status             317 non-null    object
11  HER2 status           317 non-null    object
12  Surgery_type          317 non-null    object
13  Date_of_Surgery       317 non-null    object
14  Date_of_Last_Visit    317 non-null    object
15  Patient_Status        317 non-null    object
dtypes: float64(5), object(11)
memory usage: 42.1+ KB

```

In [7]: breast_cancer_df

Out[7]:

	Patient_ID	Age	Gender	Protein1	Protein2	Protein3	Protein4	Tumour_Stage	Histology
0	TCGA-D8-A1XD	36.0	FEMALE	0.080353	0.42638	0.54715	0.273680	III	Infiltrating Ductal Carcinoma
1	TCGA-EW-A1OX	43.0	FEMALE	-0.420320	0.57807	0.61447	-0.031505	II	Mucinous Carcinoma
2	TCGA-A8-A079	69.0	FEMALE	0.213980	1.31140	-0.32747	-0.234260	III	Infiltrating Ductal Carcinoma
3	TCGA-D8-A1XR	56.0	FEMALE	0.345090	-0.21147	-0.19304	0.124270	II	Infiltrating Ductal Carcinoma
4	TCGA-BH-A0BF	56.0	FEMALE	0.221550	1.90680	0.52045	-0.311990	II	Infiltrating Ductal Carcinoma
...
329	TCGA-AN-A04A	36.0	FEMALE	0.231800	0.61804	-0.55779	-0.517350	III	Infiltrating Ductal Carcinoma
330	TCGA-A8-A085	44.0	MALE	0.732720	1.11170	-0.26952	-0.354920	II	Infiltrating Lobular Carcinoma
331	TCGA-A1-A0SG	61.0	FEMALE	-0.719470	2.54850	-0.15024	0.339680	II	Infiltrating Ductal Carcinoma
332	TCGA-A2-A0EU	79.0	FEMALE	0.479400	2.05590	-0.53136	-0.188480	I	Infiltrating Ductal Carcinoma
333	TCGA-B6-A40B	76.0	FEMALE	-0.244270	0.92556	-0.41823	-0.067848	I	Infiltrating Ductal Carcinoma

317 rows × 16 columns



```
In [8]: # Check if there are duplicate rows in the data set
breast_cancer_df.duplicated().sum()
```

Out[8]: 0

Exploratory Data Analysis

```
In [9]: breast_cancer_df.describe()
```

Out[9]:

	Age	Protein1	Protein2	Protein3	Protein4
count	317.000000	317.000000	317.000000	317.000000	317.000000
mean	58.725552	-0.027232	0.949557	-0.095104	0.006713
std	12.827374	0.543858	0.906153	0.589027	0.625965
min	29.000000	-2.144600	-0.978730	-1.627400	-2.025500
25%	49.000000	-0.350600	0.368840	-0.531360	-0.382240
50%	58.000000	0.005649	0.997130	-0.193040	0.038522
75%	67.000000	0.336260	1.612000	0.251210	0.436250
max	90.000000	1.593600	3.402200	2.193400	1.629900

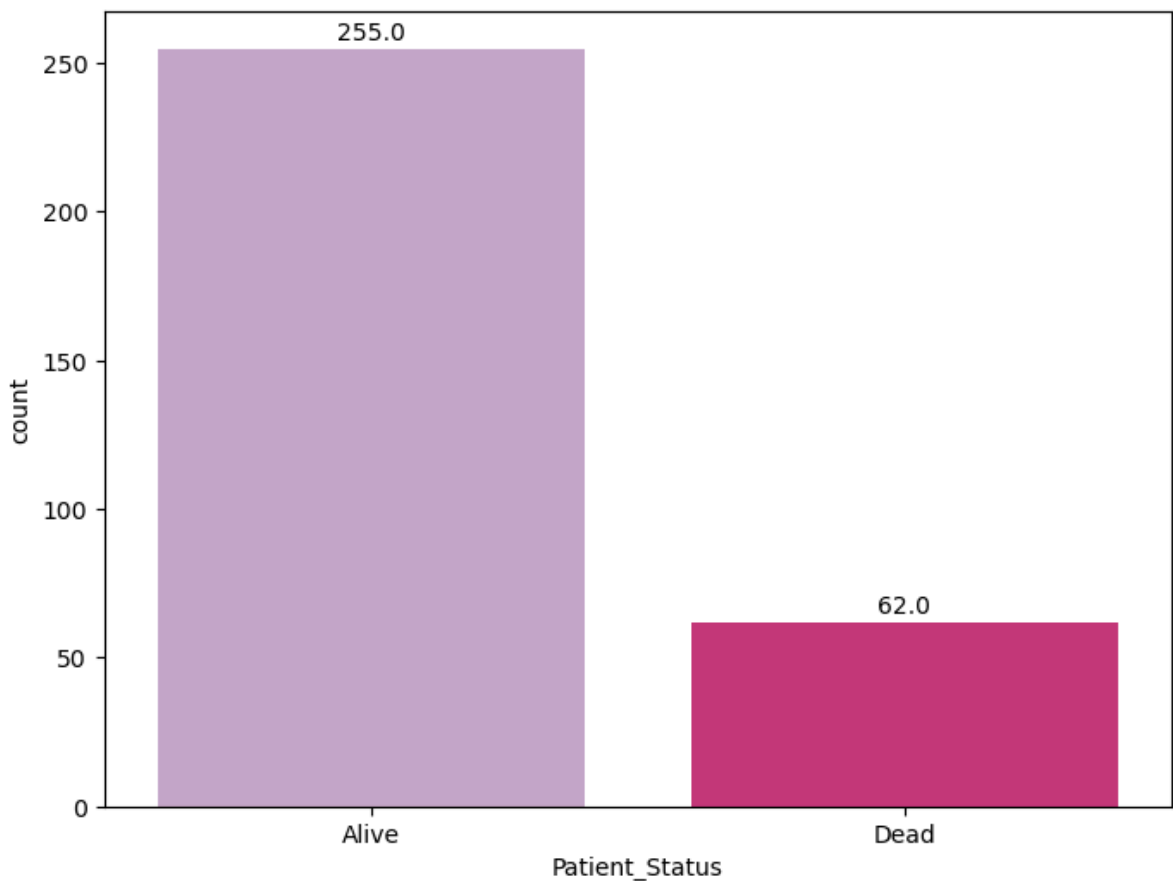
In [10]:

```
# Target distribution

# Set the figure size and create a count plot
plt.figure(figsize=(8, 6))
ax = sns.countplot(x=breast_cancer_df['Patient_Status'], palette='PuRd')

# Add labels to each bar in the plot
for p in ax.patches:
    ax.text(p.get_x()+p.get_width()/2, p.get_height()+3, f'{p.get_height()}', ha="c")

plt.show()
```




```
In [11]: breast_cancer_df['Gender'].value_counts()
```

```
Out[11]: Gender
FEMALE    313
MALE       4
Name: count, dtype: int64
```

Gender is not a useful feature in this case as there are only a few male patients

```
In [12]: breast_cancer_df['Age'].hist(bins = 50, grid = False, color='pink')
plt.xlabel(xlabel = "Age")
plt.ylabel(ylabel = "Count")
plt.title("Age Distribution of BRCA Patients")
```

```
Out[12]: Text(0.5, 1.0, 'Age Distribution of BRCA Patients')
```



Age is normally distributed, with slight right skewness and most of population is around 45-65 years old.

```
In [13]: # Separate categorical and numerical columns
categorical_columns = breast_cancer_df.select_dtypes(include=['object']).columns
numerical_columns = breast_cancer_df.select_dtypes(include=['float64', 'int64']).co
```

```
In [14]: # Calculate the Inter Quartile Range for each numerical column
Q1 = breast_cancer_df[numerical_columns].quantile(0.25)
Q3 = breast_cancer_df[numerical_columns].quantile(0.75)
```

```

IQR = Q3 - Q1

# Identify outliers for each numerical column
outliers = {}
for column in numerical_columns:
    lower_bound = Q1[column] - 1.5 * IQR[column]
    upper_bound = Q3[column] + 1.5 * IQR[column]
    outliers[column] = breast_cancer_df[(breast_cancer_df[column] < lower_bound) |

# Count the number of outliers for each numerical column
outlier_counts = {col: len(outliers[col]) for col in numerical_columns}

# Display the number of outliers for each numerical column
outlier_counts

```

Out[14]: {'Age': 0, 'Protein1': 7, 'Protein2': 0, 'Protein3': 5, 'Protein4': 4}

```


In [15]: # Remove observations with outliers from the dataframe
for col in numerical_columns:
    breast_cancer_df = breast_cancer_df[(breast_cancer_df[col] >= Q1[col] - 1.5 * I

```

```

In [16]: # Evaluate the Spearman correlation between numerical columns
correlation_matrix = breast_cancer_df[numerical_columns].corr(method='spearman')
# Plot the correlation matrix as a heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, vmin = -2, vmax = 2, cmap='PuRd')
plt.title('Correlation Matrix')
plt.show()

```

```
In [17]: # Separate the features and the target variable
X = breast_cancer_df[numerical_columns]
y = breast_cancer_df['Patient_Status']
# Encode the target variable
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(y)
```

```
In [18]: # Drop the Age column from the dataframe
breast_cancer_df = breast_cancer_df.drop('Age', axis=1)
```

```
In [19]: breast_cancer_df
```

Out[19]:

	Patient_ID	Gender	Protein1	Protein2	Protein3	Protein4	Tumour_Stage	Histology	stat
0	TCGA-D8-A1XD	FEMALE	0.080353	0.42638	0.54715	0.273680	III	Infiltrating Ductal Carcinoma	Posit
1	TCGA-EW-A1OX	FEMALE	-0.420320	0.57807	0.61447	-0.031505	II	Mucinous Carcinoma	Posit
2	TCGA-A8-A079	FEMALE	0.213980	1.31140	-0.32747	-0.234260	III	Infiltrating Ductal Carcinoma	Posit
3	TCGA-D8-A1XR	FEMALE	0.345090	-0.21147	-0.19304	0.124270	II	Infiltrating Ductal Carcinoma	Posit
4	TCGA-BH-A0BF	FEMALE	0.221550	1.90680	0.52045	-0.311990	II	Infiltrating Ductal Carcinoma	Posit
...
329	TCGA-AN-A04A	FEMALE	0.231800	0.61804	-0.55779	-0.517350	III	Infiltrating Ductal Carcinoma	Posit
330	TCGA-A8-A085	MALE	0.732720	1.11170	-0.26952	-0.354920	II	Infiltrating Lobular Carcinoma	Posit
331	TCGA-A1-A0SG	FEMALE	-0.719470	2.54850	-0.15024	0.339680	II	Infiltrating Ductal Carcinoma	Posit
332	TCGA-A2-A0EU	FEMALE	0.479400	2.05590	-0.53136	-0.188480	I	Infiltrating Ductal Carcinoma	Posit
333	TCGA-B6-A40B	FEMALE	-0.244270	0.92556	-0.41823	-0.067848	I	Infiltrating Ductal Carcinoma	Posit


301 rows × 15 columns

```
In [20]: # Plot Scatter plot on Tumor Stage vs Surgery Type
A = breast_cancer_df[breast_cancer_df.Patient_Status == "Alive"]
D = breast_cancer_df[breast_cancer_df.Patient_Status == "Dead"]

#T= breast_cancer_df[breast_cancer_df.Tumour_Stage == "III"]

plt.title("Alive vs Dead")
plt.xlabel("Tumour Stage")
plt.ylabel("Surgery Type")
plt.scatter(A.Tumour_Stage, A.Surgery_type, color = "Green", label = "Alive", alpha = 0.5)
plt.scatter(D.Tumour_Stage, D.Surgery_type, color = "Red", label = "Dead", alpha = 0.5)
```

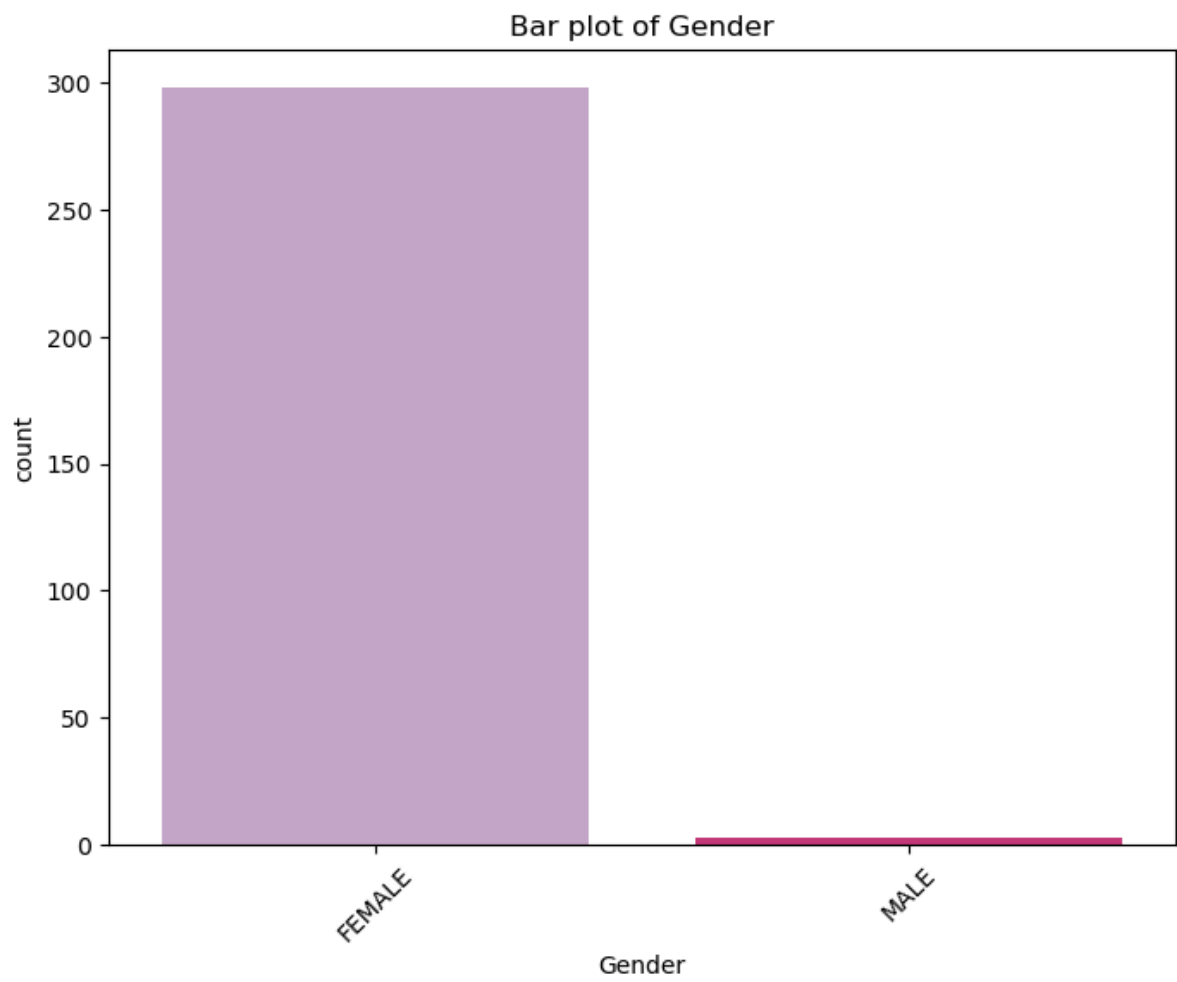
```
plt.legend()
plt.show()
```




From the above graph, we can see that for Tumor Stage 3, Lumpectomy can be performed to help the patient.


```
In [21]: # Apply kernel PCA to the numerical columns to create new features to add to the data
from sklearn.decomposition import KernelPCA
numerical_columns_pca = breast_cancer_df.select_dtypes(include=['float64', 'int64'])
kpca = KernelPCA(n_components=2, kernel='rbf', random_state=42)
kpca_features = kpca.fit_transform(breast_cancer_df[numerical_columns_pca])
breast_cancer_df['KPCA1'] = kpca_features[:, 0]
breast_cancer_df['KPCA2'] = kpca_features[:, 1]
```

```
In [22]: # Bar plot of the categorical columns to evaluate the distribution
for col in categorical_columns:
    if (col != 'Date_of_Surgery' and col != 'Date_of_Last_Visit' and col != 'Patient_Stage'):
        plt.figure(figsize=(8, 6))
        sns.countplot(data=breast_cancer_df, x=col, palette='PuRd')
        plt.title(f'Bar plot of {col}')
        plt.xticks(rotation=45)
        plt.show()
```







Bar plot of Histology



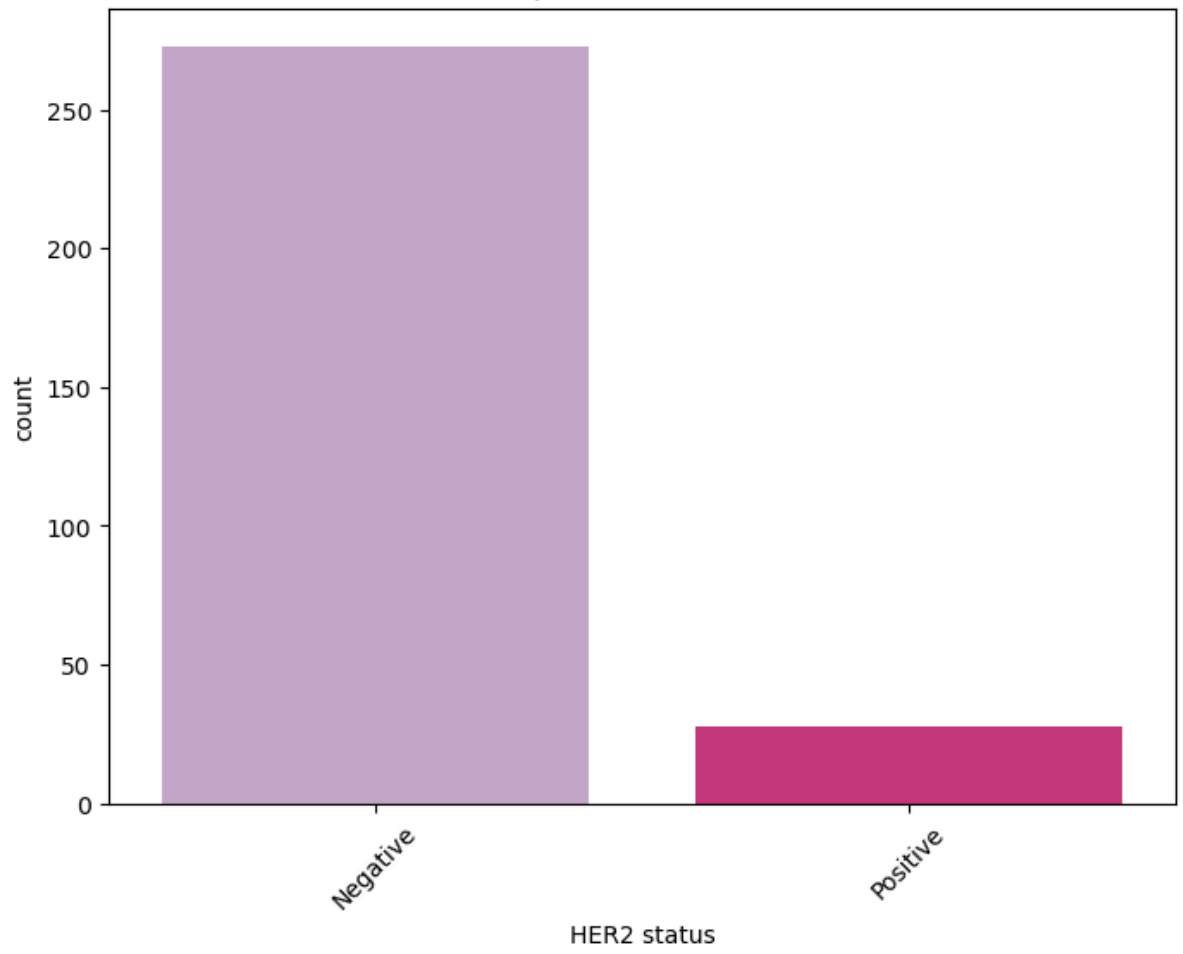
Bar plot of ER status

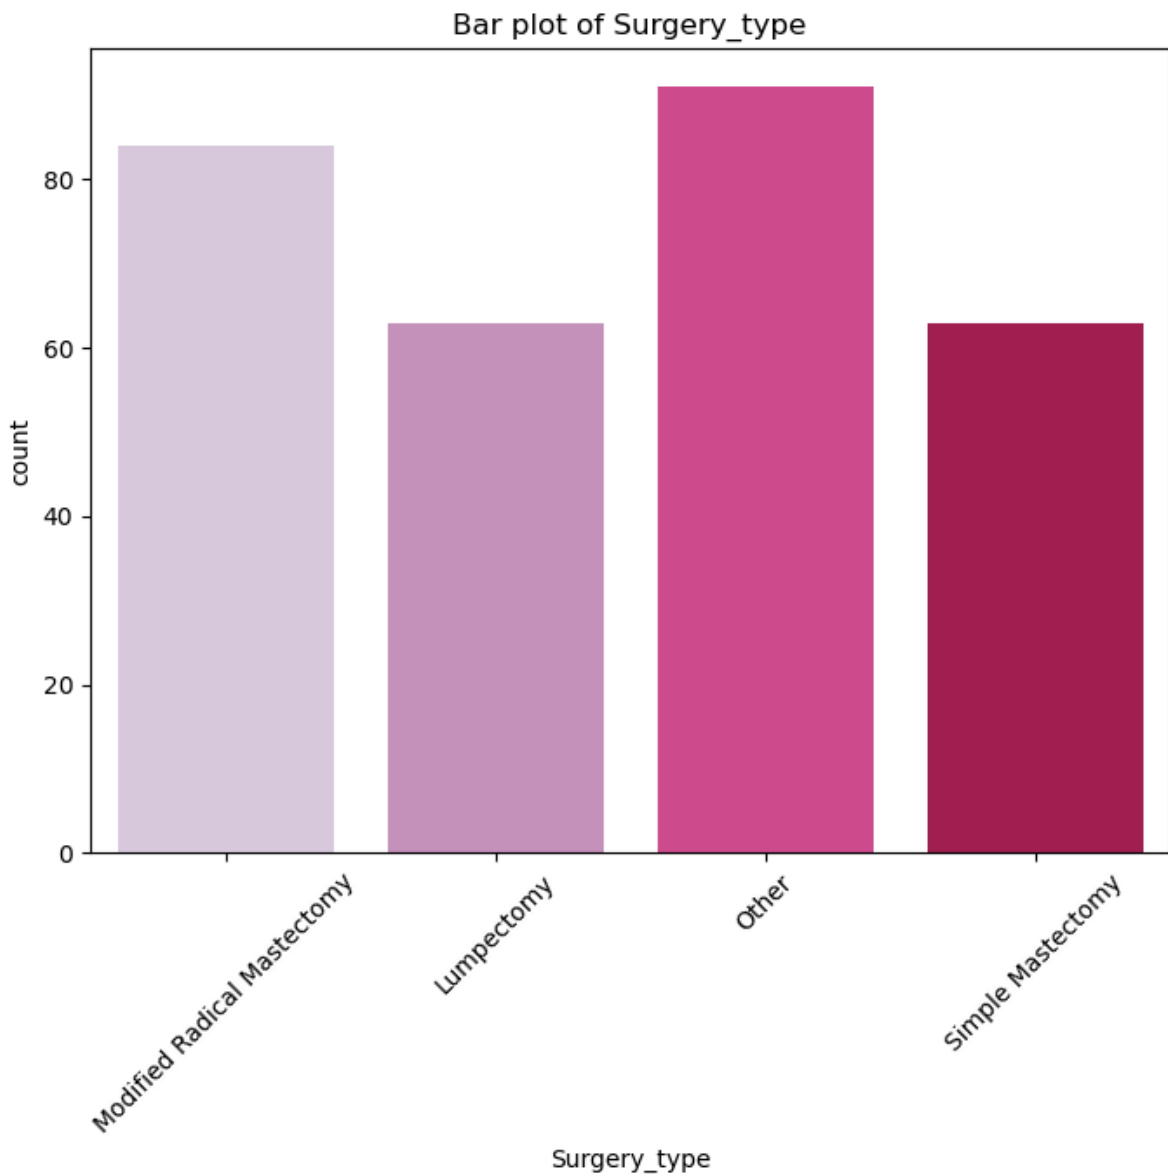


Bar plot of PR status



Bar plot of HER2 status





```
In [23]: # Drop the columns because they have one or almost unique value
breast_cancer_df=breast_cancer_df.drop(columns=['ER status','PR status','Gender'])

# Extract again the categorical columns
categorical_columns = breast_cancer_df.select_dtypes(include=['object']).columns

In [24]: # Calculate the contingency table between columns

categorical_columns = ['Tumour_Stage', 'Histology', 'HER2 status', 'Surgery_type']
for col in categorical_columns:
    contingency_table = pd.crosstab(breast_cancer_df[col], breast_cancer_df['Patient_ID'])
    print(f'Contingency Table for {col}:')
    print(contingency_table)
    print('\n')
```

Contingency Table for Tumour_Stage:

Patient_Status	Alive	Dead
Tumour_Stage		
I	51	9
II	135	33
III	56	17

Contingency Table for Histology:

Patient_Status	Alive	Dead
Histology		
Infiltrating Ductal Carcinoma	169	42
Infiltrating Lobular Carcinoma	65	14
Mucinous Carcinoma	8	3

Contingency Table for HER2 status:

Patient_Status	Alive	Dead
HER2 status		
Negative	218	55
Positive	24	4

Contingency Table for Surgery_type:

Patient_Status	Alive	Dead
Surgery_type		
Lumpectomy	55	8
Modified Radical Mastectomy	68	16
Other	68	23
Simple Mastectomy	51	12

```
In [25]: # Encode the categorical columns with one hot encoding except for Date_of_Surgery,
# Select the categorical columns to encode
categorical_columns_encode = ['Tumour_Stage', 'Histology', 'HER2 status', 'Surgery_

# Create a OneHotEncoder object
encoder = OneHotEncoder(handle_unknown='ignore', sparse_output=False, dtype=int)

# Apply encoding to the specified columns
encoded_features = encoder.fit_transform(breast_cancer_df[categorical_columns_encode

# Create a DataFrame with the encoded features
encoded_df = pd.DataFrame(encoded_features, columns=encoder.get_feature_names_out(c

# Remove the original columns and add the encoded ones
df_encoded = breast_cancer_df.drop(columns=categorical_columns_encode).reset_index(
df_encoded = pd.concat([df_encoded, encoded_df], axis=1)

# Encode the target variable Patient_Status with LabelEncoder
df_encoded['Patient_Status'] = label_encoder.fit_transform(df_encoded['Patient_Stat

df = df_encoded
```

```
In [26]: #Introduce the variable Difference_Days and divide the date of surgery and date of

# Change the data type of the 'Date_of_Surgery' and 'Date_of_Last_Visit' columns to
df['Date_of_Surgery'] = pd.to_datetime(df['Date_of_Surgery'])
df['Date_of_Last_Visit'] = pd.to_datetime(df['Date_of_Last_Visit'])

# Calculate the difference in days between the date of last visit and the date of s
df['Difference_Days'] = (df['Date_of_Last_Visit'] - df['Date_of_Surgery']).dt.days

# Extract the day, month, and year from the 'Date_of_Surgery' and 'Date_of_Last_Vis
df["Month_of_Surgery"] = df["Date_of_Surgery"].apply(lambda x : x.month)
df["Day_of_Surgery"] = df["Date_of_Surgery"].apply(lambda x : x.day)
df["Year_of_Surgery"] = df["Date_of_Surgery"].apply(lambda x : x.year)
df["Month_of_Last_visit"] = df["Date_of_Last_Visit"].apply(lambda x : x.month)
df["Day_of_Last_visit"] = df["Date_of_Last_Visit"].apply(lambda x : x.day)
df["Year_of_Last_visit"] = df["Date_of_Last_Visit"].apply(lambda x : x.year)

# Drop the 'Date_of_Surgery' and 'Date_of_Last_Visit' columns and change the displa
df = df.drop(['Date_of_Surgery', 'Date_of_Last_Visit'], axis=1)
df['Patient_Status'] = df.pop('Patient_Status')
```

Model Development

```
In [27]: df = df.drop('Patient_ID', axis=1)
X = df.drop(columns=['Patient_Status'])
y = df['Patient_Status']

# Split dataset into training and test set after feature generation
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_sta
```

```
# Check dimensions of the sets
print(f'Training set size: {X_train.shape[0]}')
print(f'Test set size: {X_test.shape[0]}')
```

Training set size: 210
Test set size: 91

```
In [28]: # Use SMOTE to balance the training set
# Create a SMOTE object
smote = SMOTE(random_state=42)

# Apply SMOTE to the training set
X_train_balanced, y_train_balanced = smote.fit_resample(X_train, y_train)

# Check the dimensions of the balanced training set
print(f'Balanced training set size: {X_train_balanced.shape[0]}')

X_train = X_train_balanced
y_train = y_train_balanced
```

Balanced training set size: 338

```
In [29]: # Function to plot a confusion matrix
def plot_confusion_matrix(y_true, y_pred, title):
    """
    Plots a confusion matrix using seaborn's heatmap visualization.

    Parameters:
    - y_true: array-like of true labels.
    - y_pred: array-like of predicted labels.
    - title: str, title of the plot.

    This function calculates the confusion matrix from true and predicted labels,
    then visualizes it with annotations and a color map of 'Blues'.
    """
    cm = confusion_matrix(y_true, y_pred) # Calculate confusion matrix
    cmd = ConfusionMatrixDisplay(cm)
    cmd.plot(cmap='PuRd');
    plt.title(title) # Set the title of the plot
    plt.xlabel('Predicted') # Set x-axis label
    plt.ylabel('Actual') # Set y-axis label

# Function to plot features importance
def plot_feature_importance(model, feature_names, title):
    """
    Plots the feature importance for tree-based machine learning models.

    Parameters:
    - model: the trained model object that has the attribute 'feature_importances_'
    - feature_names: list or array of feature names.
    - title: str, title of the plot.

    This function plots the feature importances provided by the model, sorting them
    in descending order and displaying a bar plot.
    """
    importances = model.feature_importances_
```

```

indices = np.argsort(importances[::-1])

plt.figure(figsize=(10, 6))
plt.title(title)
plt.bar(range(len(importances)), importances[indices], align='center', color='p
plt.xticks(range(len(importances)), [feature_names[i] for i in indices], rotati
plt.xlim([-1, len(importances)])
plt.xlabel('Features')
plt.ylabel('Importance')
plt.show()

```

```

In [30]: # Define the models to train, in this case, we will use XGBoost
models = {
    'XGBoost': XGBClassifier(),
    'RandomForest' : RandomForestClassifier(n_estimators=500),
    'DecisionTree' : DecisionTreeClassifier(criterion='gini',max_depth=15)
}

# Dictionary to store the accuracy for comparison
model_accuracies = {}

# Train and evaluate the models
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    model_accuracies[name] = accuracy # Store the accuracy of each model

    print(f"Model: {name}")
    print("Accuracy:", accuracy)
    print(classification_report(y_test, y_pred))

# Optional: Feature importance and confusion matrix
# Ensure these functions are defined or available in your environment
plot_feature_importance(model, X.columns, f'Feature Importances in {name}')
plot_confusion_matrix(y_test, y_pred, f'Confusion Matrix: {name}')


# Find the best model based on accuracy
best_model_name = max(model_accuracies, key=model_accuracies.get)
best_model = models[best_model_name]

```

Model: XGBoost

Accuracy: 0.7802197802197802

	precision	recall	f1-score	support
0	0.83	0.92	0.87	73
1	0.40	0.22	0.29	18
accuracy			0.78	91
macro avg	0.61	0.57	0.58	91
weighted avg	0.74	0.78	0.75	91





Model: RandomForest

Accuracy: 0.7472527472527473

	precision	recall	f1-score	support
0	0.80	0.92	0.85	73
1	0.14	0.06	0.08	18
accuracy			0.75	91
macro avg	0.47	0.49	0.47	91
weighted avg	0.67	0.75	0.70	91

Confusion Matrix: XGBoost






Model: DecisionTree


Accuracy: 0.6703296703296703

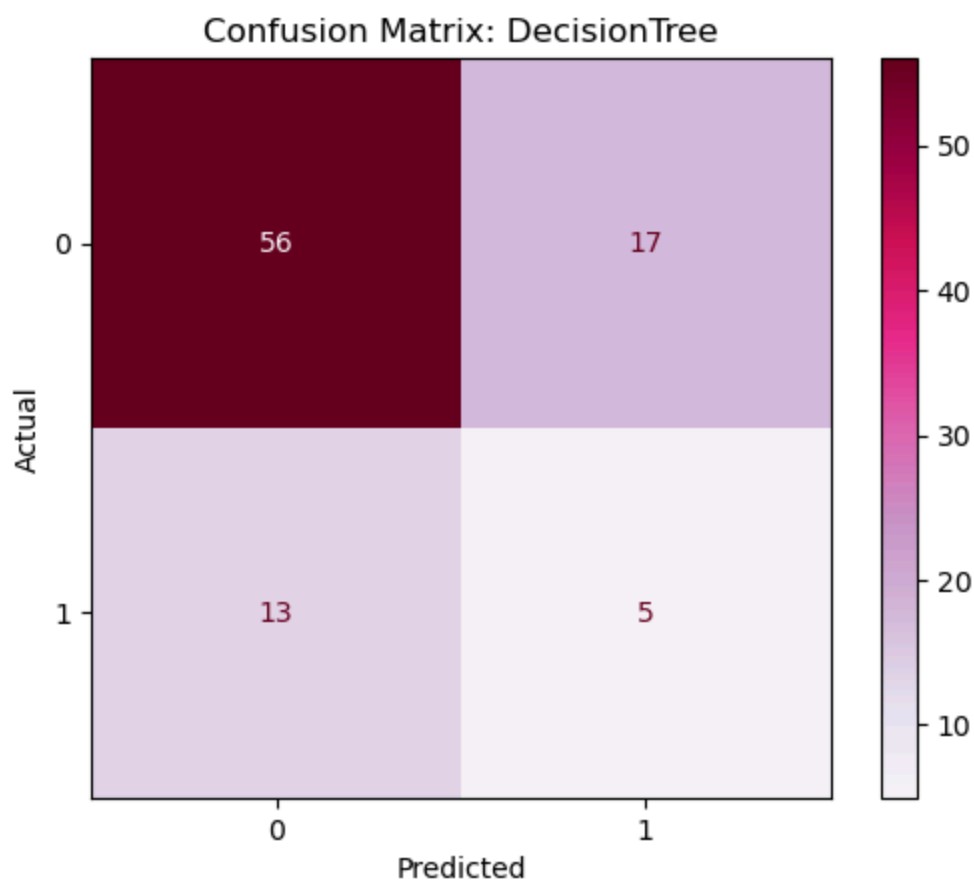
	precision	recall	f1-score	support
0	0.81	0.77	0.79	73
1	0.23	0.28	0.25	18
accuracy			0.67	91
macro avg	0.52	0.52	0.52	91
weighted avg	0.70	0.67	0.68	91

Confusion Matrix: RandomForest



Feature Importances in DecisionTree





```
In [31]: svc = SVC(probability=True)
parameters = {'kernel':('linear', 'rbf'), 'C':[1, 10]}
clf = GridSearchCV(svc, parameters, n_jobs=50)
clf.fit(X_train, y_train)
GridSearchCV(estimator=SVC(),
              param_grid={'C': [1, 100], 'kernel': ('linear', 'rbf')})
clf.best_estimator_
```

```
Out[31]: SVC
SVC(C=1, kernel='linear', probability=True)
```

Support Vector Classifier

```
In [32]: svc.fit(X_train, y_train)
y_pred = svc.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
model_accuracies[name] = accuracy # Store the accuracy of each model

print(f"Model: Support Vector Classifier")
print("Accuracy:", accuracy)
print(classification_report(y_test, y_pred))

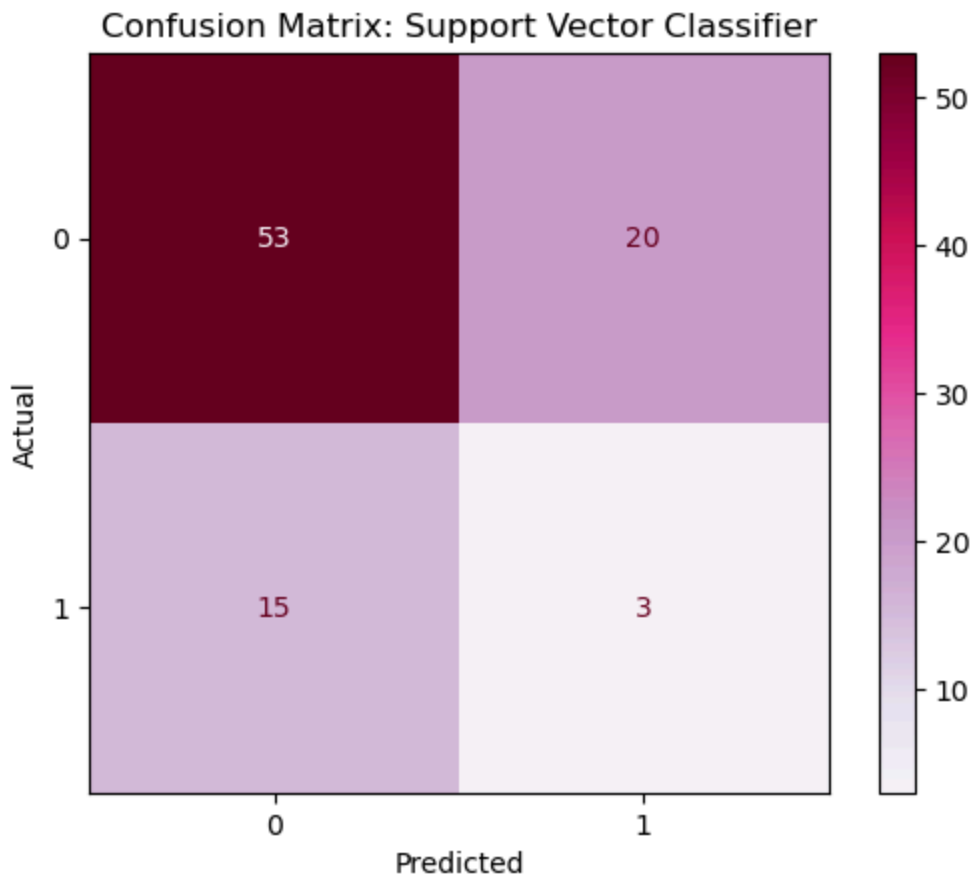
# Ensure these functions are defined or available in your environment
plot_confusion_matrix(y_test, y_pred, f'Confusion Matrix: Support Vector Classifier')
```

```
best_model_name = max(model_accuracies, key=model_accuracies.get)
best_model = models[best_model_name]
```

Model: Support Vector Classifier


Accuracy: 0.6153846153846154

	precision	recall	f1-score	support
0	0.78	0.73	0.75	73
1	0.13	0.17	0.15	18
accuracy			0.62	91
macro avg	0.45	0.45	0.45	91
weighted avg	0.65	0.62	0.63	91



Logistic Regression

```
In [33]: logit = LogisticRegression(max_iter=3500, penalty='elasticnet', solver = 'saga', l1
logit.fit(X_train, y_train)
logit_importances = logit.coef_[0]
logit_feature_importances = pd.Series(index = X_train.columns, data= np.abs(logit_i
logit_feature_importances = logit_feature_importances.sort_values(ascending=True)
logit_feature_importances.plot(kind='barh', title = 'Logistic Regression Feature I
plt.show()
```



```
In [34]: logit.fit(X_train, y_train)
y_pred = logit.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
model_accuracies[name] = accuracy # Store the accuracy of each model


print(f"Model: Logistic Regression")
print("Accuracy:", accuracy)
print(classification_report(y_test, y_pred))

# Ensure these functions are defined or available in your environment
plot_confusion_matrix(y_test, y_pred, f'Confusion Matrix: Logistic Regression')

best_model_name = max(model_accuracies, key=model_accuracies.get)
best_model = models[best_model_name]
```

```
Model: Logistic Regression
Accuracy: 0.6263736263736264
```

	precision	recall	f1-score	support
0	0.84	0.66	0.74	73
1	0.26	0.50	0.35	18
accuracy			0.63	91
macro avg	0.55	0.58	0.54	91
weighted avg	0.73	0.63	0.66	91



```
In [35]: print("Best Model: ", best_model)
```

```
Best Model: XGBClassifier(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=None, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None, feature_types=None,
    gamma=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=None, max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=None, max_leaves=None,
    min_child_weight=None, missing=nan, monotone_constraints=None,
    multi_strategy=None, n_estimators=None, n_jobs=None,
    num_parallel_tree=None, random_state=None, ...)
```